

Scaling Data-Race Freedom Analysis with Array Projections

Paul Maynard Tiago Cogumbreiro

University of Massachusetts Boston

SERPL 2023

Overview

- ▶ Data-Race Freedom in GPU kernels.
- ▶ Array projections – nonlinear expressions.
- ▶ Limitations of SMT solvers.
- ▶ Translating expressions to multidimensional accesses.
- ▶ Opens up analysis of a broader variety of programs.


GPU programming

- ▶ Massively parallel execution (thousands of threads).
- ▶ Blocks of threads.
- ▶ Single Instruction, Multiple Threads.
- ▶ Thread-global memory and thread-local memory.

GPU programming model

```
void scal(float a, float *x)
{
    int i = blockIdx*blockDim + threadIdx;
    x[i] = a*x[i];
}
```

- ▶ Each thread runs this function with the same arguments.
- ▶ `threadIdx` is different for each thread



What is a data-race?

- ▶ Concurrent programs with shared memory can have data-races.
- ▶ Occur when two threads access the same memory location at the same time.
- ▶ At least one thread is writing.
 - ▶ One write and one write – data-race.
 - ▶ One write and one read – data-race.
 - ▶ One read and one read – no data-race.

Data-Race Freedom (DRF)

To prove a program is data-race free:

- ▶ Examine all concurrent memory accesses.
- ▶ Any accesses between synchronization points.
- ▶ Prove that all pairs of accesses are safe with each other.

Checking DRF with SMT

1. Find set of concurrent accesses.
2. For each 2 elements in set, can indices be the same?
3. Translate accesses into an SMT (Satisfiability Modulo Theories) formula.
 - ▶ If formula satisfiable, then data-race.
 - ▶ If formula unsatisfiable, then DRF.

Example – checking DRF

```
void scal(float a, float *x)
{
    int i = blockIdx*blockDim + threadIdx;
    x[i] = a*x[i]; // read x[i], write x[i]
}
```

$$\text{tid}_1 \neq \text{tid}_2 \wedge \text{bid} \cdot \text{bdim} + \text{tid}_1 = \text{bid} \cdot \text{bdim} + \text{tid}_2$$

- ▶ *Thread-local variables are instantiated* for 2 symbolic threads. [FSE'10]
- ▶ `bid` and `bdim` are constant across threads, so this is solvable.
- ▶ No possible values that could satisfy this formula.
- ▶ Therefore, it is DRF.

Challenge – Nonlinearity

Projections of multidimensional accesses

- ▶ A more complicated problem:

```
void kernel(float* paths, int S, int T) {  
    int step = gridDim * blockDim;  
    for (int i = threadIdx; i < S ; i += step) {  
        for (int t = 0; t < T ; t++) {  
            paths[i + S * t] = f(t);  
                ^^^^^^^^^  
        }  
    }  
}
```

$$i_1 < S \wedge i_2 < S \wedge i_1 + S \cdot t_1 = i_2 + S \cdot t_2$$

- ▶ This time the access involves thread local variables.

Nonlinear formulas

$$i_1 < S \wedge i_2 < S \wedge i_1 + S \cdot t_1 = i_2 + S \cdot t_2$$

- ▶ This time the variables are thread local.
- ▶ Undecidability.
- ▶ SMT solvers can't solve all nonlinear expressions.

Nonlinear formulas


$$i_1 < S \wedge i_2 < S \wedge i_1 + S \cdot t_1 = i_2 + S \cdot t_2$$

- ▶ This time the variables are thread local.
- ▶ Undecidability.
- ▶ SMT solvers can't solve all nonlinear expressions.

Array projections


Multidimensional array representation

- Many nonlinear accesses are actually projections of 2d arrays.



Multidimensional array representation

- ▶ Many nonlinear accesses are actually projections of 2d arrays.



- ▶ Our contribution: To make these solvable by SMT solvers, we can reverse this transformation.

Rewriting accesses

```
void kernel(float paths[][], int S, int T) {  
    int step = gridDim * blockDim;  
    for (int i = threadIdx; i < S ; i += step) {  
        for (int t = 0; t < T ; t++) {  
            paths[t][i] = f(t);  
            ^^^^  
        }  
    }  
}
```

$$i_1 < S_1 \wedge i_2 < S_2 \wedge i_1 = i_2 \wedge t_1 = t_2$$

- This formula is now analyzable.

Rewriting projections

How do we rewrite?

1. Identifying patterns (potentially obfuscated or nested).

`dimension * x_index + y_index`

2. Identifying potential rewrites.
3. Determining which rewrites are sound.
4. Choosing the appropriate rewrite.

$A[a * b + a] \implies A[a][c]$

or

$A[a * b + a] \implies A[b][c]?$

Identifying projections

- ▶ One option: hard-coded patterns.
- ▶ Doesn't catch all instances.

- ▶ For example:

$$100 * (a * b + c) + 1$$

- ▶ How often do nonstandard instances occur in practice?

Rewrite rules

- ▶ Rewrite expressions using series of rules
- ▶ Put all expressions into a normal form.
- ▶ Identify projections from normal form.
- ▶ Example:

$$100 * (a * b + c) + 1 \\ \implies 100 * a * b + 100$$

Egraphs

- ▶ Equivalence-graphs.
- ▶ Track equivalence classes of expressions.
- ▶ Allow fast and efficient rewriting of terms. [egg'21]

Soundness and ambiguity

When is rewriting sound?

- ▶ Need to make sure we don't remove data-races.

- ▶ Example:

$$A[a * b + c] \implies A[a][c]$$

when

$$a = 1$$

$$b = 5$$

$$c = 6$$

$$A[1 * 5 + 6] \implies A[1][6]$$

$$A[11]$$

$$A[2 * 5 + 1] \implies a[2][1]$$

- ▶ Must guarantee that $c < b$
 - ▶ Dimensions must be same across all accesses.

Ambiguous expressions

$100 * a * b + 100 * c$

- ▶ What is the index and what is the dimension?
 - ▶ 6 possible interpretations
- ▶ Heuristics:
 - ▶ Prefer thread globals and constants to be dimension.
 - ▶ Prefer loop variables to be indices.
 - ▶ Soundness rules can help choose:
 - ▶ Must have same dimension across all accesses
 - ▶ Y-index must be less than the dimension.
- ▶ Could send multiple interpretations to the solver.

Open questions/Future work

- ▶ Testing/analysis on GPUVerify Cav14 dataset.
 - ▶ What patterns exist, and in what proportion?
 - ▶ What level of reweriting/searching is needed?
- ▶ Can this fix be automatically applied to most programs?
- ▶ Is it sound, or are there data-races it eliminates?
- ▶ Use in additional contexts
 - ▶ Code repair/synthesis.
 - ▶ Precondition inference.

Conclusion

- ▶ Data-races in GPU programming.
- ▶ Array projections.
- ▶ Rewriting projections back to multidimensional accesses.
 - ▶ Recognizing & rewriting patterns.
 - ▶ Soundness of rewrites.
 - ▶ Handling ambiguity.